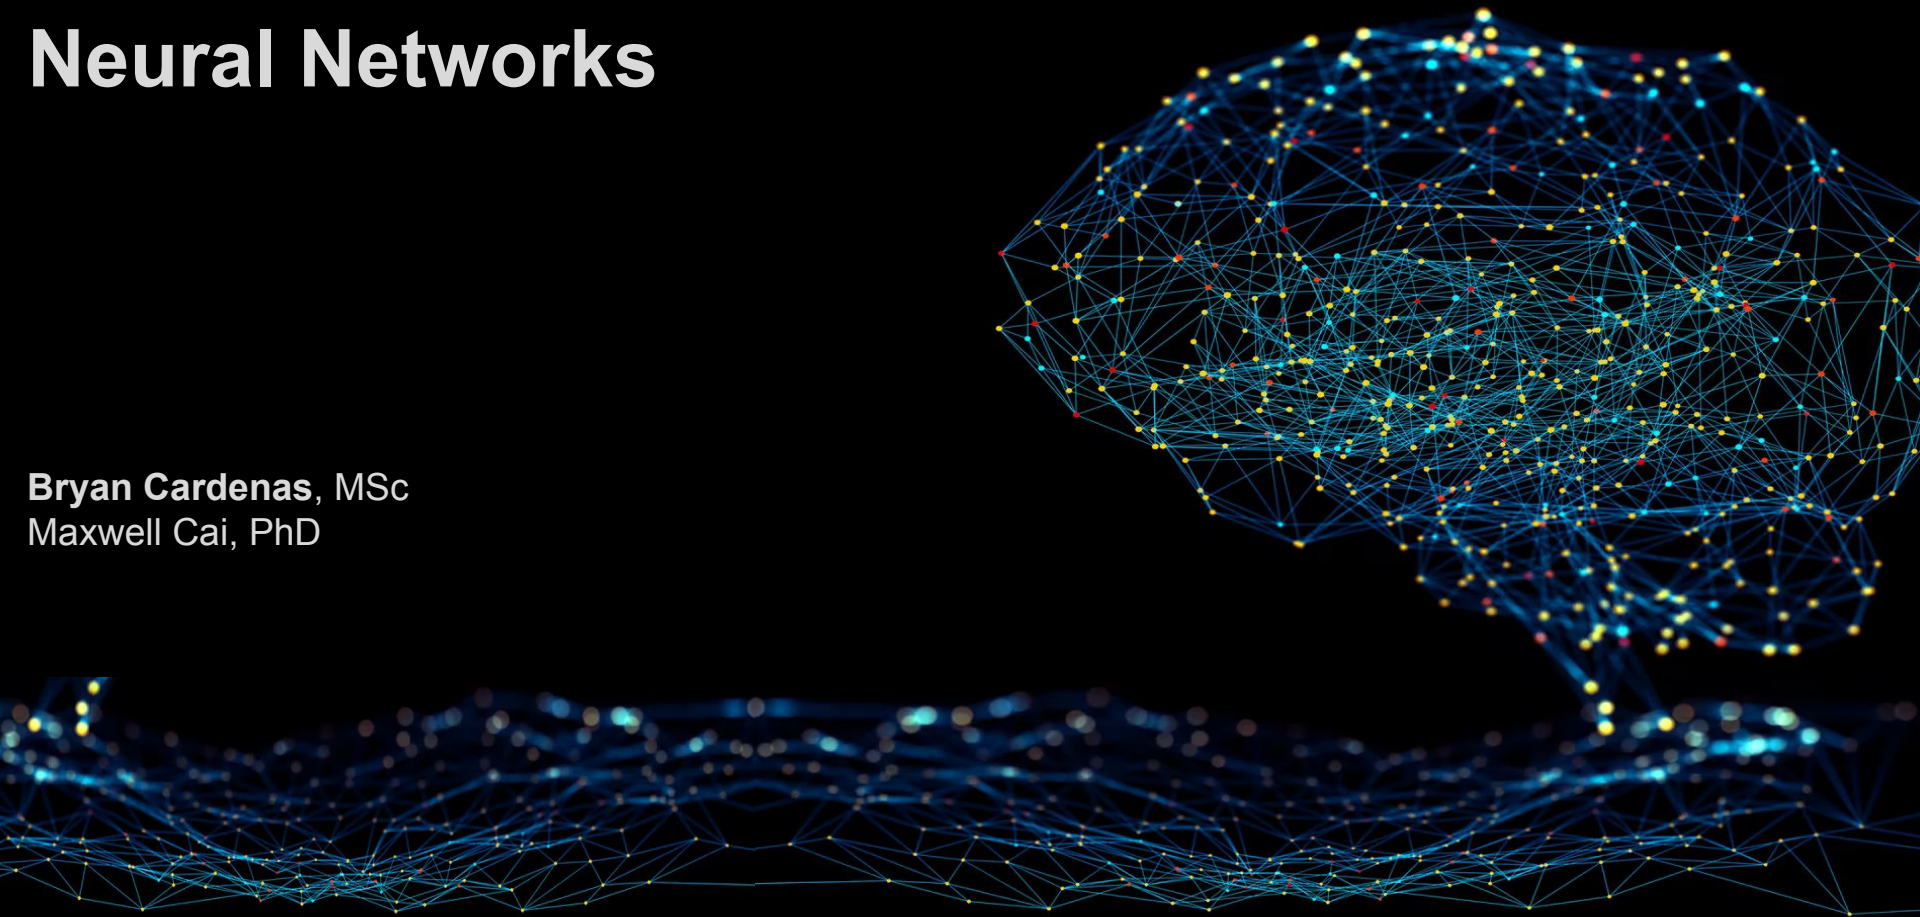


Neural Networks

Bryan Cardenas, MSc
Maxwell Cai, PhD



Background Prerequisites

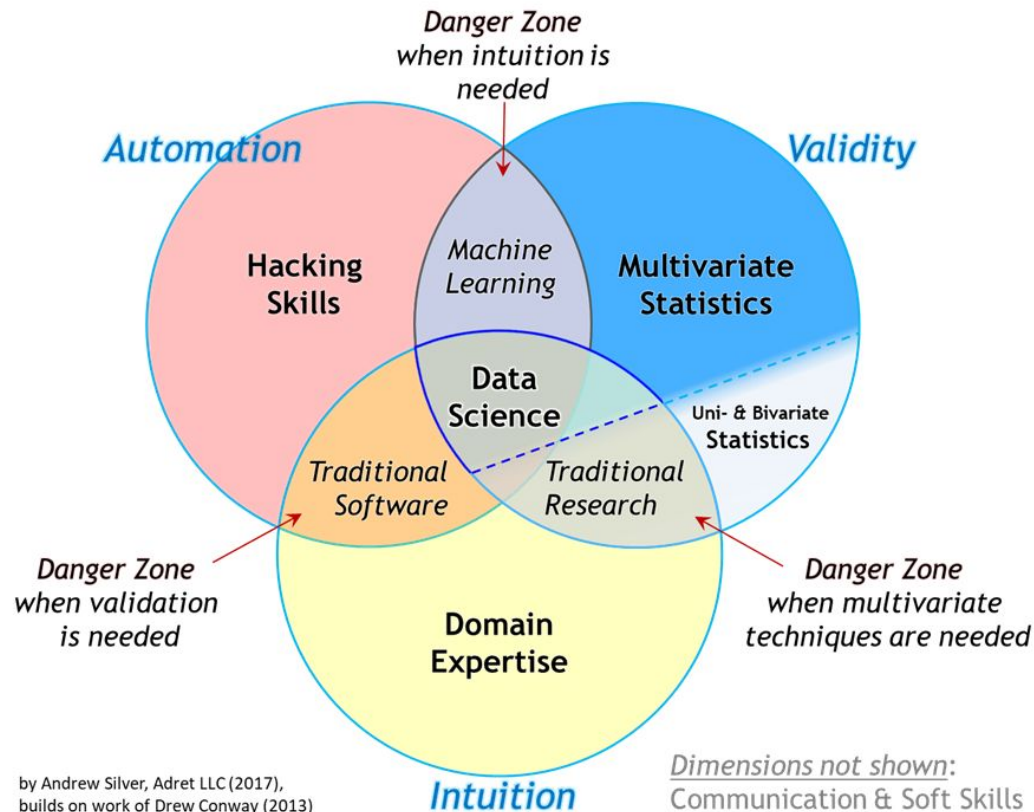
Programming

R / Python

Statistics, Calculus

Machine Learning / Deep Learning

Parallel Computing



Course Plan

1. General Introduction to ML and DL, basic principles
2. Algorithms and Models
3. Convolutional Neural Networks
4. Generative Models
5. Recurrent Neural Networks
6. (brief) Reinforcement Learning
7. Afternoon: High Performance Machine Learning

What is Machine Learning?

It is **NOT**:

Mimicking human intelligence

Robotics

Deep Learning




***ML** is the study of computer [algorithms](#) that can improve automatically through experience and by the use of data.^[1] It is seen as a part of [artificial intelligence](#).*

- *wikipedia*

AI vs ML vs DL

Artificial Intelligence

Having computers to exert
Intelligent behaviour



AI vs ML vs DL


Artificial Intelligence

Having computers to exert
Intelligent behaviour



Machine Learning

Perform tasks without
Explicitly programmed
from data



AI vs ML vs DL


Artificial Intelligence

Having computers to exert
Intelligent behaviour




Machine Learning

Perform tasks without
Explicitly programmed
from data



Deep Learning

Use (deep) neural networks



AI vs ML vs DL


Artificial Intelligence

Having computers to exert
Intelligent behaviour




Machine Learning

Perform tasks without
Explicitly programmed
from data



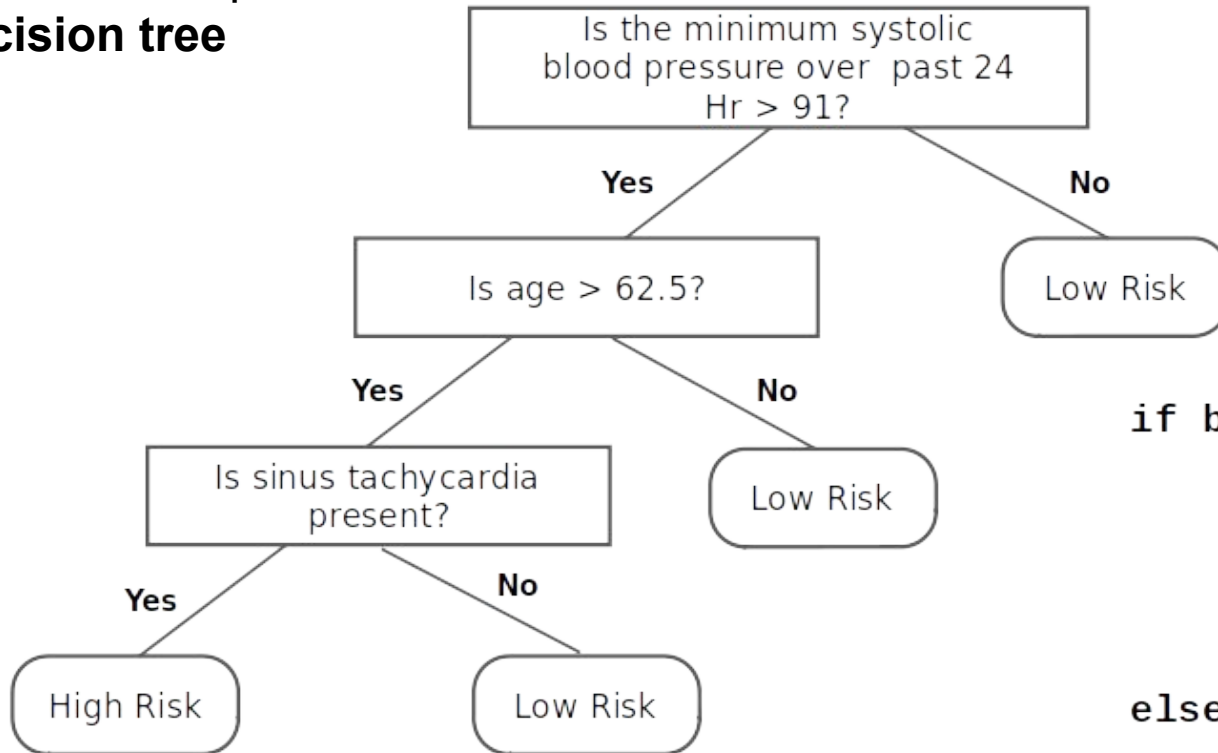
Deep Learning

Use (deep) neural networks



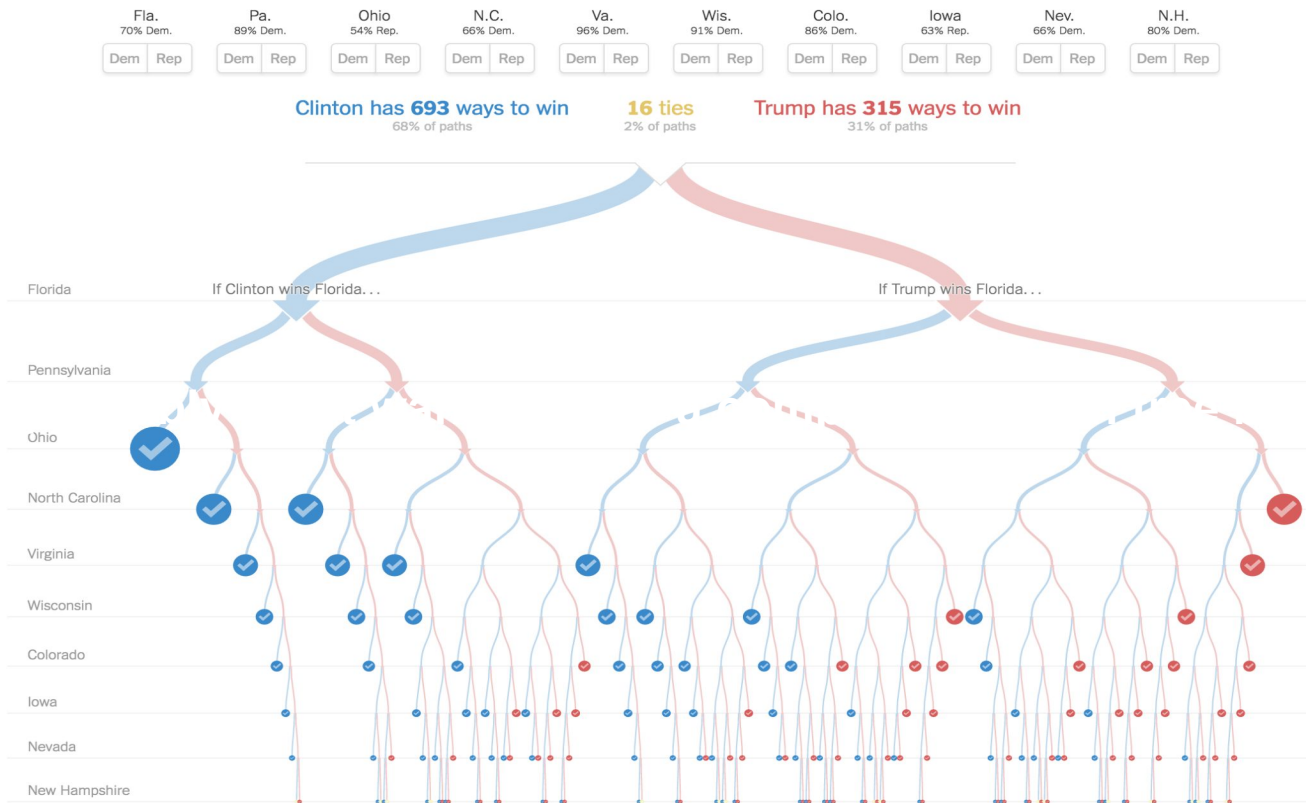
Why Machine Learning?

Think of a simple
decision tree

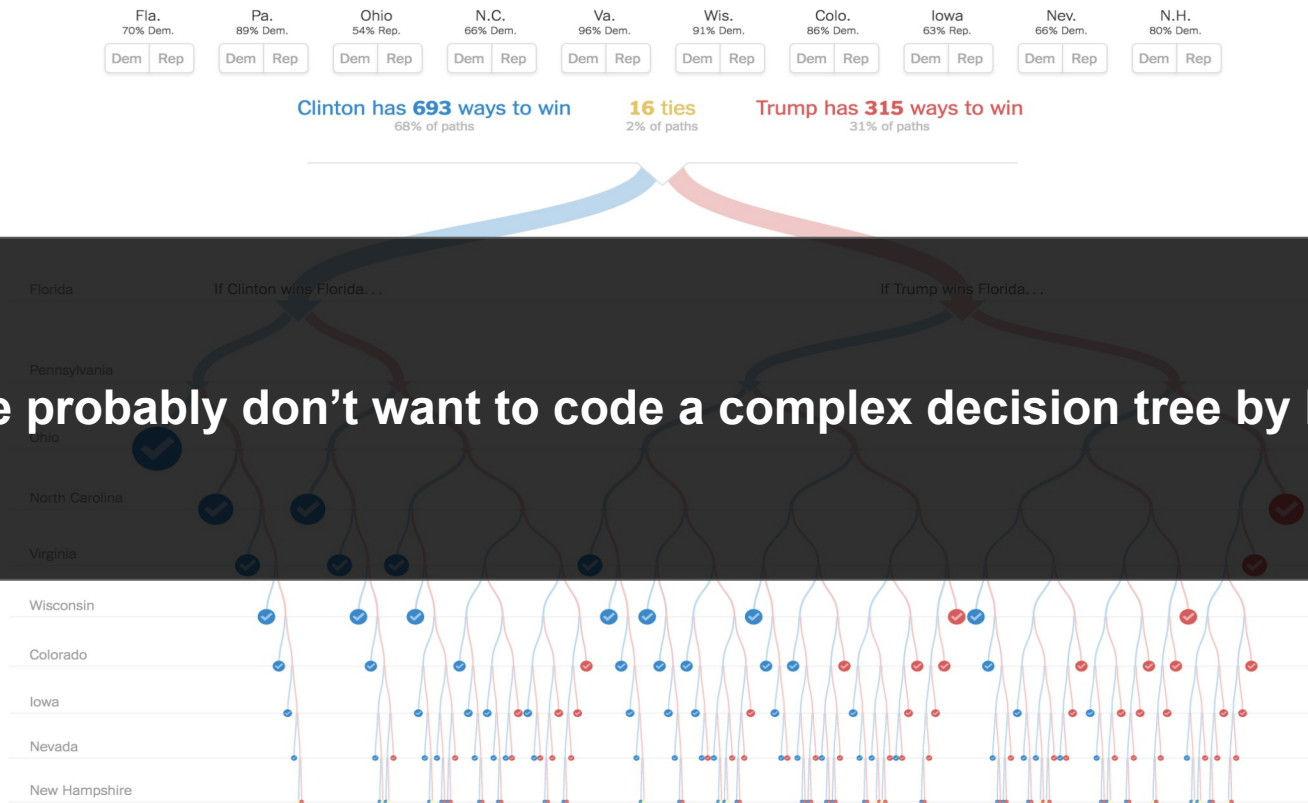


```
if blood_pressure > 91:
    if age > 62.5:
        if sinus_tach:
            ...
        else:
            ...
    else:
        ...
```

Why Machine Learning?



Why Machine Learning?



Why Machine Learning?

What is a **dog**?



Why Machine Learning?

What is a **dog**?

Driven by Data




Categories of Machine Learning

Supervised

Learn from labels


Regression, Classification



Unsupervised

Detect patterns in the data

Clustering, Dimensionality reduction




Reinforcement

Learn from mistakes

Control, Gaming




Neural Networks



Use a (deep) neural network to **approximate** an unknown **function**

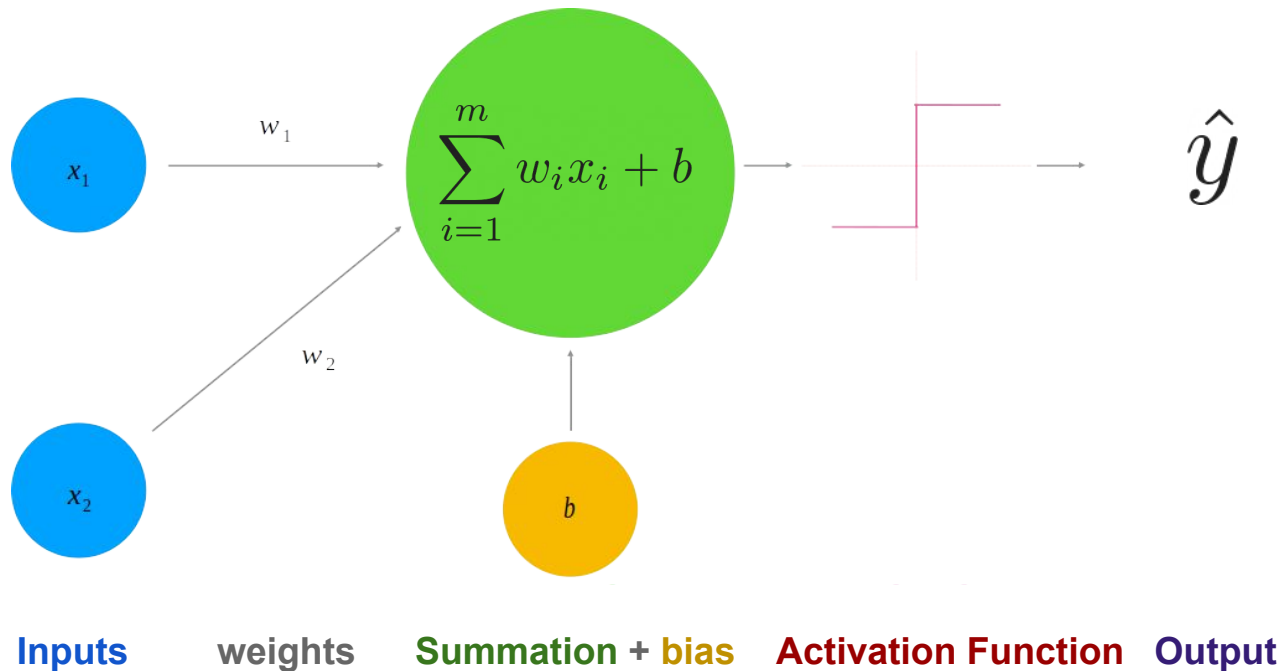
Neural Networks



$$\hat{y} = f_{NN}(x_1, x_2, \dots, x_n)$$


Use a (deep) neural network to **approximate** an unknown **function**

Anatomy of Neural Networks





Perceptron & Activation

Binary Classification Task





Perceptron & Activation

Binary Classification Task






Perceptron & Activation

Binary Classification Task






Perceptron & Activation

Binary Classification Task






Perceptron & Activation

Binary Classification Task



Limitation of Linear Single-Layer Classifiers





XOR Problem

Possible solutions:


- Add more layers (deep learning)
- Map into another (higher dimensional) space

Kernel Trick




Adding Layers

House price prediction



Adding Layers


House price prediction



During the **optimization** process,
The machine learns to **encode** a
representation that maps the
input to the output


Adding Layers

A deep neural network **encodes** the **representation** in an increasingly abstract way



Adding Layers

- Neural networks are made from neurons and edges
- A collection of neurons in a layer
- The output of previous layer is used as an input to the next layer
- The input layer is data input and the output is a prediction
- Anything in between is **hidden**
- Layers are represented as vectors
- Edges are usually represented as matrices - **The weights**
- We train the weights



Adding Layers

Universal Approximation Theorem

“Given a neural network with a **single hidden layer** of **sufficient size**, the network can Approximate any continuous function”

In other words:

- There exists a true function relating the inputs to the outputs
- A neural network can approximate this function to arbitrary precision given sufficient layer size
- The required layer size can be extremely large and grow rapidly with the dimensionality of the problem


Use **multiple hidden layers** —————> Encoding becomes increasingly more abstract

Adding Layers

Estimate $\hat{y} = f_{NN}(x_1, x_2, \dots, x_n)$

Loss

Ground Truth




Adding Layers

Estimate $\hat{y} = f_{NN}(x_1, x_2, \dots, x_n)$

Loss $L(y, \hat{y}) = L(W, b) = (y_i - \hat{y}_i)^2$

Ground Truth




Adding Layers

Estimate $\hat{y} = f_{NN}(x_1, x_2, \dots, x_n)$

Loss
$$L(y, \hat{y}) = L(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Ground Truth




Adding Layers


Estimate $\hat{y} = f_{NN}(x_1, x_2, \dots, x_n)$

Loss $L(y, \hat{y}) = L(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$


Ground Truth $\mathbf{X} = (x_1, \dots, x_m), y$



Training a Neural Network




Training a Neural Network




$$L(y, \hat{y}) = L(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Training a Neural Network



$$L(y, \hat{y}) = L(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Training a Neural Network




$$L(y, \hat{y}) = L(W, b) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Learning Rate


$$\mathbf{w}_{j+1} = \mathbf{w}_j - \alpha \nabla L(\mathbf{w}_j, b)$$

Overfitting


Do we want the lowest loss? **Not Really**




Underfitted



Good Fit/Robust



Overfitted




Overfitting

We have to combat overfitting


A few ways to do so is by:

- Simply stopping training earlier
- Dropout: deactivate a neuron and its connections for the forward propagation with a certain probability
- Decay the value of your weights over time




Training: Backwards Propagation

Forwards Propagation



Training: Backwards Propagation

Backwards Propagation




A (deep) neural network is a deeply **nested functions**:

- We need to compute the gradient for each layer
- Apply the **chain rule**
- This is **backpropagation**

Training: Backwards Propagation

Backwards Propagation




A (deep) neural network is a deeply **nested functions**:

- We need to compute the gradient for each layer
- Apply the **chain rule**
- This is **backpropagation**

Training: Backwards Propagation

Backwards Propagation




A (deep) neural network is a deeply **nested functions**:

- We need to compute the gradient for each layer
- Apply the **chain rule**
- This is **backpropagation**

Training: Backwards Propagation

$$\frac{\partial L(y, \hat{y})}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \mathbf{a}_i} \frac{\partial \mathbf{a}_i}{\partial \mathbf{w}_{i,j}}$$

Backwards Propagation




A (deep) neural network is a deeply **nested functions**:

- We need to compute the gradient for each layer
- Apply the **chain rule**
- This is **backpropagation**

Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Many more! We can design our own!

Commonly used loss functions

Regression


- Mean Squared Error (MSE)
- Mean Squared Log Error
- Mean Absolute Error

Binary Classification

- Binary cross-entropy
- Hinge Loss

Multi-Class Classification


- Multi-class cross-entropy
- Kullback-Leibler Divergence



Cross-entropy loss outputs a log probability



Optimizers

In what way should we change the weights?



Optimizers

In what way should we change the weights?




General Workflow of ML


You need to know your data and your models well.

Artificial Intelligence still heavily relies on human intelligence


Imbalanced Training set




Undersampling



Oversampling




Data Normalization




A process to transform the input **data** in a **well-behaved** form

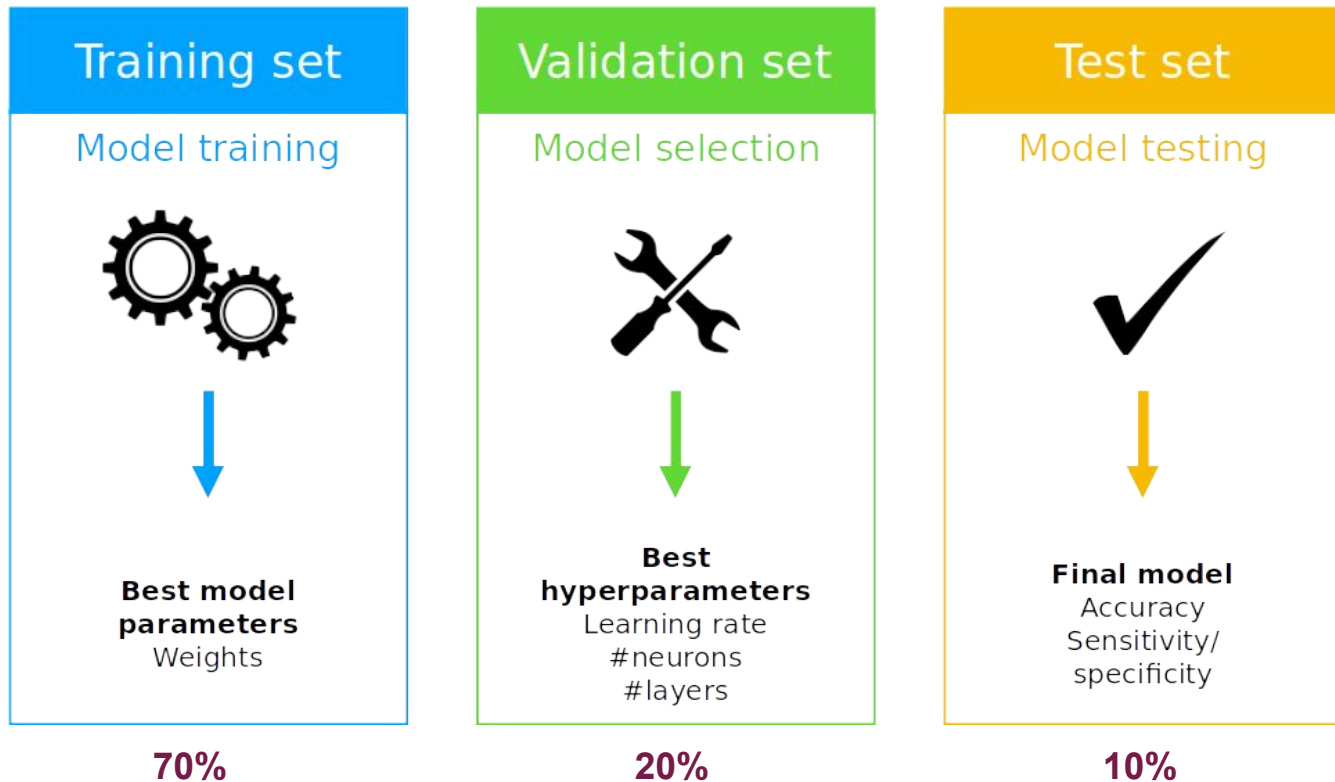
Ratings per movie



Log ratings per movie



Dataset Splitting




Network Evaluation

		Prediction outcome		total
		p	n	
actual value	p'	True Positive	False Negative	P'
	n'	False Positive	True Negative	N'
total		P	N	

Accuracy = (TP + TN) / (TP + FN + FP + TN)

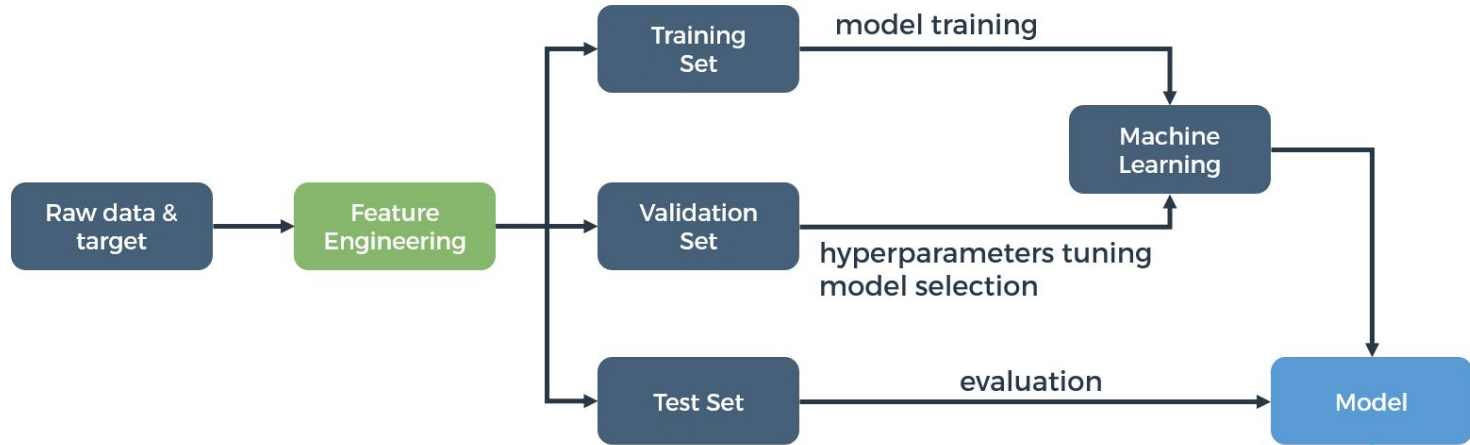
Precision (p) = TP / (TP + FP)

Recall (r) = TP / (TP + FN)



Workflow

TRAINING




PREDICTING



DL Frameworks

In DL you need to

- Define neurons and layers
- Define loss function
- Calculate losses
- Calculate gradient
- Propagate backward
- Update weights
- Existing frameworks exist:
 - TensorFlow (Keras)
 - Torch
 - Jax
 - MXNet



Open Datasets

Datasets

Find and use datasets or complete tasks. [Learn more.](#)

[+ New Dataset](#)

Help the community by creating and solving Tasks on datasets!



Search 29,853 datasets

Feedback Filter

PUBLIC

Sort by: Hottest



Hotel booking demand

Jesse Mostipak

19 days 1 MB 10.0 1 File (CSV) 1 Task

270



Big Five Personality Test

Bojan Tunguz

14 days 159 MB 9.7 3 Files (CSV, other)

134



Startup Investments (Crunchbase)

Andy_M

14 days 3 MB 8.8 1 File (CSV)

92

Open Tasks

Can we predict the possibility of a bo...

0 Submissions · In Hotel booking demand

Visualize US Accidents Dataset

12 Submissions · In US Accidents (3.0 million...

What to watch on Netflix ?

4 Submissions · In Netflix Movies and TV Sh...

The state that has the highest number...

5 Submissions · In US Accidents (3.0 million r...

**Processed, balanced,
well-behaved and labelled
datasets**

tensorflow.org/datasets

kaggle.com/datasets

topepo.github.io/caret/data-sets.html

[github.com/awesomedata/awesome-pu
blic-datasets](https://github.com/awesomedata/awesome-public-datasets)

Take Home Messages

Machine Learning

New paradigm of programming, driven by data

An optimization process

Deep Learning

A subfield of ML


Relies on deep neural networks

Learns to encode the input data using many layers of concept hierarchies

Take Home Messages

In a neuron:

- ... the main job is to calculate a weighted average
- ... the decision is made through the activation function



In a neural network:

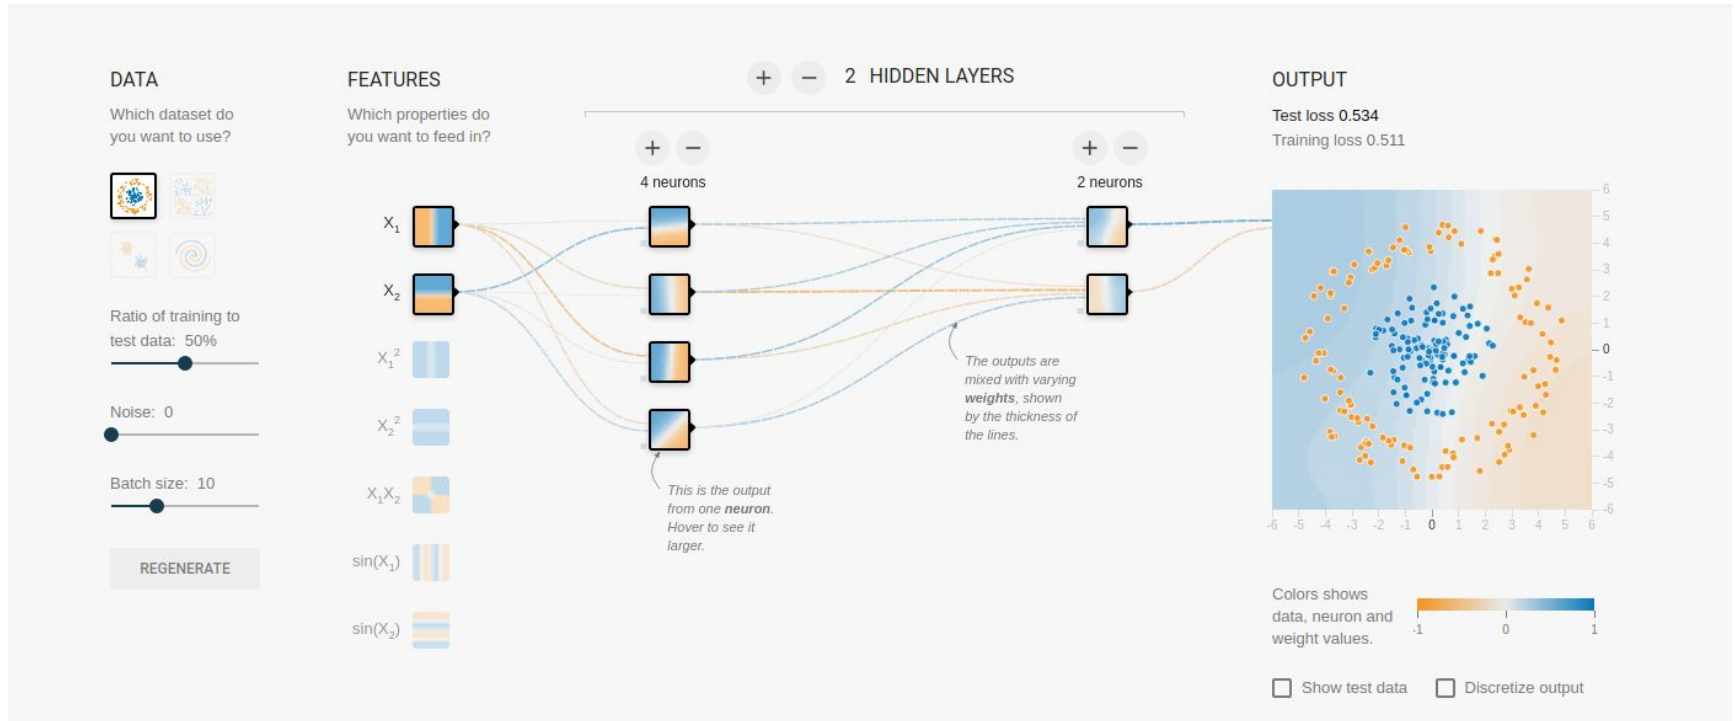
- ... losses are calculated using the loss function
- ... losses are calculated by comparing the labels and the prediction
- ... predictions are made through forward propagation
- ... weights are updated through the backward propagation process
- ... optimizers are used to decide the weights updating strategies

In a deep learning workflow:

- ... the heavy lifting is mostly done by DL frameworks
- ... open datasets are crucial for benchmarking and bootstrapping DNNs



Live Demo



<https://playground.tensorflow.org/>